

Test Dummy

405-555-0100 | randommail6969@gmail.com | Springfield, IL

PROFESSIONAL SUMMARY

Applied **AI/ML** Analyst with four years of experience building production data pipelines, developing scalable transformations in **PySpark** and **SQL**, and architecting cloud infrastructure on **AWS**. Proven track record in asset management and financial services environments, with hands-on expertise in **Snowflake** optimization, **Apache Airflow** orchestration, and enterprise data governance. Skilled at collaborating with data science and ML teams to design infrastructure that supports model training, deployment, and performance monitoring in production.

TECHNICAL SKILLS

Languages & Processing: Python, SQL (complex joins, window functions, query optimization), PySpark, Scala, Bash

Cloud & AWS Services: S3, IAM, KMS, Glue, Lambda, EMR, Step Functions, Kinesis, MSK, MWAA, Athena

Data Warehousing & Governance: Snowflake (micro-partitioning, clustering, warehouse sizing, query tuning, data sharing), Redshift, data lineage, encryption, least-privilege access, deduplication, auditability

Orchestration & ML Infrastructure: Apache Airflow, AWS Step Functions, batch and streaming pipeline design, feature pipeline development, model performance monitoring

Data Modeling: Dimensional and star schemas, normalized models, banking data concepts, enterprise analytics patterns

WORK EXPERIENCE

Applied AI/ML Analyst

Jan 2025 – Present

Example Corp | Remote

- Build and maintain end-to-end batch and streaming data pipelines (ingestion → transformation → curation → serving) on **AWS**, leveraging **S3**, **Glue**, **EMR**, **Kinesis**, and **Step Functions** to move financial data from source systems into analytics-ready layers serving 12+ downstream teams.
- Develop scalable transformations using **PySpark** and **SQL**, optimizing for performance and cost; enabled Adaptive Query Execution, tuned executor memory to 48GB, and resolved data skew via salting on join keys, reducing average job runtime by 38% on the largest daily batch while cutting cluster costs by 26%.
- Implement and optimize **Snowflake** solutions including schemas, tables, views, micro-partitioning and clustering strategies, warehouse sizing, and query tuning across high-frequency analytical workloads; introduced materialized views and optimized clustering keys, reducing warehouse credit consumption by 22% over one quarter.
- Design and maintain feature pipelines for **ML** scoring models, implementing statistical drift detection and schema validation checks to ensure training and inference environment parity; integrated pipeline outputs with model performance tracking to flag degradation against defined thresholds.
- Set up **Airflow**/**MWAA** DAGs with retry logic, alerting, and operational runbooks; 94% of production jobs now run within SLA without manual intervention, down from 67% baseline.
- Enforce **KMS** encryption, least-privilege **IAM** policies, audit logging, and data retention rules across all pipelines to meet governance standards required in financial services contexts.
- Continuously evaluate emerging **AI** technologies to optimize **ETL** workflows; piloted LLM-assisted **SQL** code generation for ad-hoc pipeline development, reducing build time by 35% for the analytics team.

Data Engineer (Contract)

Jan 2024 – Dec 2024

Sample Industries LLC | Springfield, IL

- Designed dimensional and star schema data models for an enterprise analytics platform supporting regulatory reporting and dashboards across 15+ internal business units, applying standard banking data concepts and governance patterns throughout.
- Built **Snowflake** data loading pipelines with access controls and data sharing patterns; redesigned clustering keys and introduced materialized views, which dropped warehouse credit consumption by 22% over the quarter.
- Wrote **Python** automation scripts for ingestion, validation, and reconciliation across 50+ source feeds, enforcing completeness, timeliness, accuracy, and deduplication checks at each stage of the pipeline.
- Set up **CI/CD** workflows in Git with unit and integration tests as part of code review standards; the team saw a 31% reduction in production defects after the first full sprint under the new process.
- Handled incident triage and root-cause analysis for **Airflow**-orchestrated pipelines, then wrote operational runbooks that became the team's standard reference for on-call response.
- Ran daily data quality reconciliation between source systems and **Snowflake** targets, flagging discrepancies before they reached downstream consumers.
- Collaborated with risk and compliance stakeholders to document data lineage, define retention policies for regulated datasets, and ensure auditability across the platform.

Data Engineering Analyst

Aug 2021 – Dec 2022

Example Systems | Remote

- Wrote complex **SQL** queries with window functions, multi-table joins, and aggregations for large-scale financial domain reporting; rewrote several legacy stored procedures that had been running slow for years, cutting query time by up to 45%.
- Built batch data pipelines on **AWS** using **S3**, **Glue**, and **Lambda** to ingest and transform structured and semi-structured source files, with validation and reconciliation controls baked into each stage.
- Implemented security and governance controls including encryption at rest, least-privilege access policies, and audit logging to satisfy compliance requirements on banking client engagements.
- Collaborated with technology, risk, and compliance teams to document data lineage, define retention schedules, and deliver datasets that met both regulatory and business requirements.
- Contributed to Git-based **CI/CD** pipeline setup and participated in code reviews across a 12-person engineering squad, which helped standardize how the team shipped and tested pipeline changes.
- Ran **PySpark** jobs on **EMR** clusters processing multi-hundred-GB daily files; worked with senior engineers to understand partition tuning, broadcast join patterns, and dynamic allocation strategies.

Data Engineering Intern

Dec 2020 – May 2021

Example Systems | Remote

- Wrote **SQL** transformation scripts and **Python** utilities to support **ETL** pipeline development for enterprise analytics projects, handling ingestion and data quality checks.
- Ran completeness and accuracy validation on source-to-target mappings and contributed reconciliation reports reviewed by senior engineers before production sign-off.
- Gained hands-on experience with **AWS S3** and **Lambda**, and ran **Spark** jobs under mentorship, building foundational knowledge for distributed processing patterns.

EDUCATION

Master's Degree, Computer Science, University of Alabama, Tuscaloosa, AL (GPA: 3.10/4.0)

Bachelor's Degree, Computer Science, GITAM, Visakhapatnam, India (GPA: 3.00/4.0)